

# Digital Design with Verilog

RISC-V : Instruction Formats

Lecture 28 : RISC-V Part 3





# Disclaimer

I do not own all the slides which I am going to present.  
The content is mostly from

- GIAN course titled “**Next-Generation Semiconductors: RISC-V, AI, TL-Verilog**” by Steeve Hoover.  
<https://github.com/silicon-vlsi/gian-course-2024-IITBBS>
- Introduction to Computer Architecture by Hasan Baig,  
<https://www.hasanbaig.com>
  - Most of the slides are copied/directly from his course content, please visit his site for reference.
- Digital Design and Computer Architecture, RISC-V Edition by Sarah L. Harris and David Money Harris.
- Other online publicly available sources. I tried my best to acknowledge the source wherever possible. **I too might have missed some sources too 😊**



# How to represent Instructions ?

- Everything in computer is represented with binary and so the **Instructions** as well
  - Called machine code
- How many bits should we use to encode instructions?
- Are we using the same number of bits to encode all instructions?
  - Do all instructions have the same length?

```
add          rd, rs1, rs2
addi         rd, rs1, immd
lw           rd, offset(rs1)
sw           rs2, offset(rs1)
beq          rs1, rs2, LABEL
```

# How to represent Instructions ?

## Design Principle 4: Good design demands good compromises

Keep formats as similar as possible

instruction

32 bits


- RISC-V instructions
  - Encoded as 32-bit instruction words
  - Small number of formats, encoding operation code (opcode), register numbers, ...
- Regularity!

# RISC-V Instruction Format

- The layout of bits in instruction words is **instruction format**
  - How do we use 32-bit bits to specify operation code (opcode), registers, immediate, offset, etc? How many bits for each?
- RISC-V has six instruction formats
  - R, I, S, SB, U, and UJ

# R-Format

# R-Format



- For instructions that have 3 registers as operands
- Fields in R-type
  - opcode: 7-bit operation code
  - rd: destination register number
  - rs1: first source register number
  - rs2: second source register number
  - funct3: additional function code
  - funct7: even more function code




# R-Format



a = b + c

add x9,x20,x21



0000000\_10101\_10100\_000\_01001\_0110011<sub>2</sub>


0000 0001 0101 1010 0000 0100 1011 0011<sub>2</sub>

015A04B3<sub>16</sub>





# R-Format



*Any instruction like:*

**instruction          rd, rs1, rs2**

*can be represented with R-Format.*

## R-Format: Knowledge Check


	Funct7
ADD	0b 000 0000
SUB	0b 010 0000

- What RISC-V instruction does this represent?  
Choose from one of the four options below.
  - The numbers are in decimal

funct7	rs2	rs1	funct3	rd	opcode
32	9	10	0	11	51

- A. sub      x9,      x10,      x11
- B. add      x11,      x9,      x10
- C. sub      x11,      x10,      x9
- D. sub      x11, x9, x10


# R-Format



	type	opcode	funct3	funct7
add	R	0110011	000	0000000
sub	R	0110011	000	0100000
sll	R	0110011	001	0000000
slt	R	0110011	010	0000000
sltu	R	0110011	011	0000000
xor	R	0110011	100	0000000
srl	R	0110011	101	0000000
sra	R	0110011	101	0100000
or	R	0110011	110	0000000
and	R	0110011	111	0000000


**I-Format**

# I-Format



- Immediate arithmetic and load instructions
  - rs1: source or base address register number
  - immediate: constant operand, or offset added to base address
    - 2's-complement, sign extended
- Since only 12 bits are kept in machine code, the immediate must be in  $[-2^{11}, +2^{11} - 1]$  or  $[-2048, 2047]$

# I-Format



addi x1, x2, 32



When executing the instruction, processor builds a 32-bit immediate **imm** by extending the sign bit




# I-Format

immediate	rs1	funct3	rd	opcode
12 bits	5 bits	3 bits	5 bits	7 bits
	type	opcode		funct3
addi	I	0010011		000
slli	I	0010011		001
slti	I	0010011		010
sltiu	I	0010011		011
xori	I	0010011		100
srli	I	0010011		101
srai	I	0010011		101
ori	I	0010011		110
andi	I	0010011		111

In **slli**, **srli**, and **srai**, only lower 5 bits of the immediate are used for shift amount. 5 bits are enough for 32-bit registers!


## I-Format Instructions: srli and srai



**srli** x1, x2, 16

More than 31 times shift is not needed, therefore the maximum immediate value, i.e. 31, can be represented in 5 bits.


Remaining 7 bits can be repurposed as an additional funct7 field









# I-Format Instructions: srli and srai



srli x1, x2, 16



srai x1, x2, 16



srli	I	0010011	101	0000000	← funct7
srai	I	0010011	101	0100000	



# I-Format Instructions: load

- Load instructions

lw rd, offset(rs1)

A. R-format

B. I-format

R:	funct7	rs2	rs1	funct3	rd	opcode
I:	imm [11:0]		rs1	funct3	rd	opcode



# I-Format Instructions: load

imm[11:0]	rs1	funct3	rd	opcode
12 bits	5 bits	3 bits	5 bits	7 bits

lw x1, 32(x2)

0000 0010 0000	00010	010	00001	000 0011
----------------	-------	-----	-------	----------

	type	opcode	funct3
lb	I	0000011	000
lh	I	0000011	001
lw	I	0000011	010
ld	I	0000011	011
lbu	I	0000011	100
lhu	I	0000011	101
lwu	I	0000011	110

Do you see the pattern in funct3?

How about store?



# RISC-V Instruction Formats

- Store instructions have rs2, but not rd


sw rs2, offset(rs1)

Can any of the following two formats be used for store instructions?


R:	funct7	rs2	rs1	funct3	rd	opcode
I:	imm [11:0]		rs1	funct3	rd	opcode

# S-Format

# S-Format



- Different immediate format for store instructions
  - rs1: base address register number
  - rs2: source operand register number
  - immediate: offset added to base address
    - Split so that rs1 and rs2 fields always remain in the same place



# R/I/S-type Instructions Summary

Instruction	Format	funct7	rs2	rs1	funct3	rd	opcode
add (add)	R	0000000	reg	reg	000	reg	0110011
sub (sub)	R	0100000	reg	reg	000	reg	0110011

7 bits    5 bits    5 bits    3 bits    5 bits    7 bits

Instruction	Format	immediate	rs1	funct3	rd	opcode
addi (add immediate)	I	constant	reg	000	reg	0010011
lw (load word)	I	address	reg	010	reg	0000011

12 bits    5 bits    3 bits    5 bits    7 bits

Instruction	Format	immed-iate	rs2	rs1	funct3	immed-iate	opcode
sw (store word)	S	address	reg	reg	010	address	0100011

7 bits    5 bits    5 bits    3 bits    5 bits    7 bits

# U-Format



# RISC-V Instruction Format

	31	27	26	25	24	20	19	15	14	12	11	7	6	0
R	funct7				rs2		rs1		funct3		rd		Opcode	
I	imm[11:0]						rs1		funct3		rd		Opcode	
S	imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode	

**U-Format**



# U-Format

- Recall Load 32-bit 0x003D0500 into register x19
  - LUI set the higher 20 bits (adjusted for sign bit from the lower 12 bits)
  - ADDI set the lower 12 bits
    - Immediate is sign extended

```
lui x19, 0x3D0 # 0x003D0
```

x19: 


0000 0000 0011 1101 0000	0000 0000 0000
--------------------------	----------------

```
addi x19, x19, 0x500
```

x19: 

0000 0000 0011 1101 0000	0101 0000 0000
--------------------------	----------------

# U-Format



- Fields in U-type

- opcode: operation code
- rd: destination register number
- imm[31:12]:

The higher 20 bits of the 32-bit immediate **the processor gets**

- Opcode:

LUI: **0b\_011\_0111**

## Example:

**lui**      **x1, 0xABCDE**      # the 32-bit immediate is 0xABCDE000




# U-Format

# the immediate written in assembly language

```
lui    x1, 0xABCDE
```

The immediate field in the machine code



The 32-bit immediate the processor uses is 0xABCDE000



# **SB-Format**

# RISC-C Instruction Format

R:	funct7	rs2	rs1	funct3	rd	opcode
I:	imm [11:0]		rs1	funct3	rd	opcode
S:	imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode
U:	imm [31:12]				rd	opcode

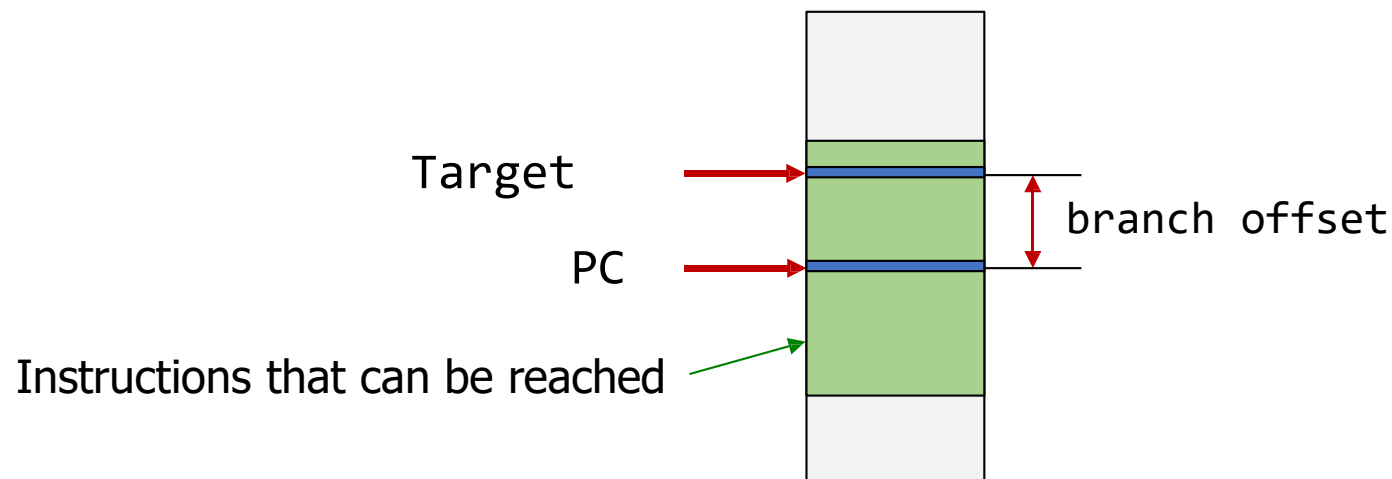
beq rs1, rs2, label

All branches are similar.

- Which format is the best for branches?
- What problem do we have?

# Offset in Branches

- The full instruction address is 32 bits
  - There are only 12 bits to encode the target address
- Instead of the full address, we store **the offset**, a smaller number in the 12 bits
  - The offset is the distance between the target address and PC
  - The target address is relative to PC (**PC-relative addressing**)







# Offset in Branches

- When encoding branches, assembler calculates the offset:

$$\text{branch offset} = \text{target address} - \text{PC}$$

> 0:      jump forward  
== 0:     the branch itself  
< 0:      jump backward


- When processor executes the instruction, it
  - Forms a 32-bit immediate from the 12 bits in machine code
  - Calculates the target address as

$$\text{target address} = \text{PC} + \text{imm}$$

$$\text{imm} == \text{branch offset}$$


# Recalling Address Alignment

- Words and half-words have alignment issue
- When aligned,
  - Addresses of words must be a multiple of 4
  - Addresses of half-words must be a multiple of 2
- RISC-V does not require data to be aligned
- However, instructions must be aligned



# Offset Address Boundary

- Each instruction is 32 bits wide  
→ offset of 4 bytes
- Sometimes the instructions are 16-bits wide (in C extension)  
→ Can be located at any even addresses  
→ offset of 2 bytes
- RISC-V keeps the ability of branching to any even addresses
- Because of the insignificance of right most bit (i.e., always zero), the Immediate data field can be represented in **13 bits**
- Instead of -2048 to 2046, we can now access -4096 to 4094 in multiples of 2



# SB-Format

Instruction	Format	immed-iate	rs2	rs1	funct3	immed-iate	opcode
sw (store word)	S	address	reg	reg	010	address	0100011

7 bits


5 bits

5 bits

3 bits

5 bits

7 bits



11

10

9

8

7

6

5

4

3

2

1

0

12-bits Immediate data field

Always 0  
=> multiple of 2

Branch instructions	Imm[12]	Imm[10:5]	rs2	rs1	fn3	Imm[4:1]	Imm[11]	opcode
---------------------	---------	-----------	-----	-----	-----	----------	---------	--------

Branch instructions	Immediate field		rs2	rs1	fn3	Immediate field		opcode
---------------------	-----------------	--	-----	-----	-----	-----------------	--	--------

7 bits

5

5

3

5 bits

7 bits

bits

bits

bits



# SB-Format

**bne x10, x11, 2000** // if x10 != x11, goto location address 2000

2000 = 0 0111 1101 0000

0	111 110	01011	01010	001	1000	0	1100011
Imm[12]	Imm[10:5]	rs2	rs1	fn3	Imm[4:1]	Imm[11]	opcode

Branch ranges for 13 bits = -  $2^{12}$  to  $2^{12} - 1$  = -4096 to 4094

Represent -4096 offset in the immediate format of branch instruction



# RISC-V Instruction Summary

Simple subset that shows most aspects

Arithmetic/logical: `add`, `sub`, `and`, `or`

Memory reference: `lw`, `sw`

Branch: `beq`

Type	Instruction	Opcode	Funct3	Funct7
R-Type	<code>add</code>	011 0011	000	000 0000
R-Type	<code>sub</code>	011 0011	000	010 0000
R-Type	<code>and</code>	011 0011	111	000 0000
R-Type	<code>or</code>	011 0011	110	000 0000
I-Type	<code>lw</code>	000 0011		
S-Type	<code>sw</code>	010 0011		
SB-Type	<code>beq</code>	110 0011		

# RISC-V Instruction Summary

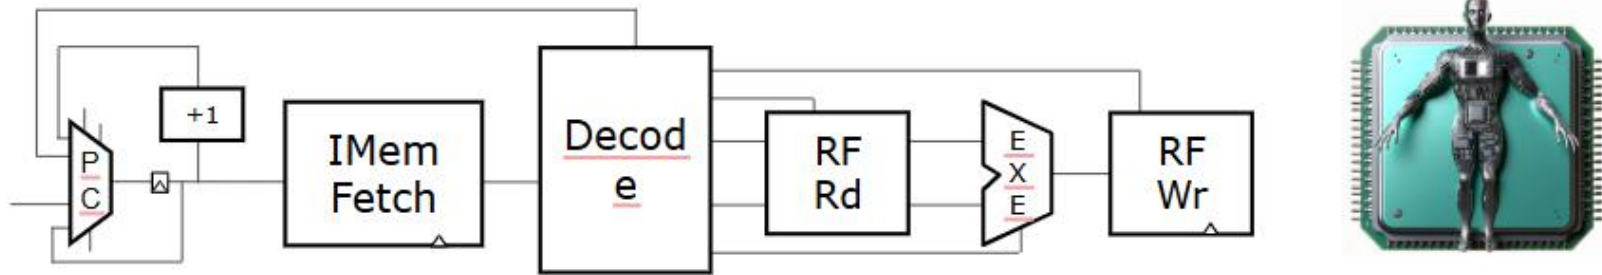
Fields are at the same location in all encoding formats

opcode, funct3, rs1, rs2, rd

Name (Bit position)	Fields					
	31:25	24:20	19:15	14:12	11:7	6:0
(a) R-type	funct7	rs2	rs1	funct3	rd	opcode
(b) I-type	immediate[11:0]		rs1	funct3	rd	opcode
(c) S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode
(d) SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode

How are the instructions executed?

# Human CPU



1. Step through the execution of the program, one instruction at a time.
  - **Next PC:** Adjust PC to the next instruction (sequential or branch target) (initially to 0)
  - **Fetch:** Identify the binary instruction at the PC
  - **Decode:** Flip the Instruction Memory card to find the decoded (assembly) instruction
  - **RF Read:** "Get" the source register values from the register file
  - **Execute:** Determine the result of the operation
  - **RF Write:** Update (in pencil) the value in the destination register
  - Repeat
2. What does the program do?
3. Clean up.



**Thank you**